

A Why model-level alignment is not enough

model-in-institution evaluation

model-level evaluation stops here

Aligned agents

each passes
model-level tests

Shared resource

memory, compute,
code, budgets

Emergent failure

depletion, collusion,
duplicated work

Institution

human-designed rules,
monitoring, repair

Outcomes + costs

cooperation, burden,
false punishment

Evaluate the model-in- institution, not the model alone.

Individually aligned agents can still fail as a group. Institutions can prevent that, but they add costs and failures of their own, so the outcomes are hypotheses to test, not results.

B What an institution can use

66 coded mechanisms in 18 themes and six families (Fig. 4)

