


A Why model-level alignment is not enough

model-in-institution evaluation

model-level evaluation stops here




Evaluate the model-in-institution, not the model alone.

Individually aligned agents can still fail as a group. Institutions can prevent that, but they add costs and failures of their own, so the outcomes are hypotheses to test, not results.

B What an institution can use

66 coded mechanisms in 18 themes and six families (Fig. 4)



LLM-agent evidence: ● tested ⊙ proposed only ○ none yet