


FROM ALIGNED MODELS TO GOVERNED AGENT SOCIETIES

Aligned in isolation




An individually aligned agent avoids harm — alone.

Sharing a commons



...but collectively unsafe:
depletion, gaming, cascades.

Wrapped in an institution



Institutions sustain
cooperation.

POSITION

Multi-agent AI safety needs institutional design. Even individually aligned agents can produce collectively unsafe outcomes when sharing finite resources, so we should evaluate the model-in-institution, not the model alone.