







A. Familiar commons examples

in human social dilemmas



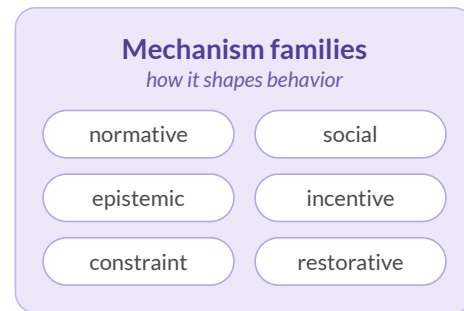
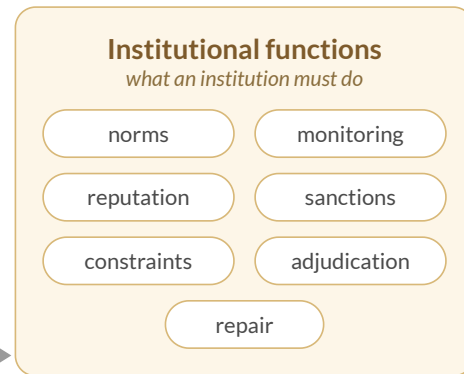
B. Artificial commons examples

in LLM agent systems



C. Design space

institutional functions × mechanisms



Multi-agent AI safety requires intentional institutional design. Even individually aligned agents can produce collectively unsafe outcomes when sharing finite resources. We believe agent systems should be studied as governed systems whose institutions shape collective behavior.

— our central claim

Contribution 1

Aligned individuals are not enough; institutions are needed.

Contribution 2

Distinguishing 7 functions from 6 mechanism families.

Contribution 3

An evaluation agenda for cooperation-shaping institutions.